

## Task 12: Building a Multilayer Perceptron (MLP) using TensorFlow and Keras

**Objective:** In this assignment, you will learn to build, train, and evaluate Multilayer Perceptron (MLP) models using sklearn, TensorFlow and Keras. You will work with two datasets—the Iris datasets—and experiment with different MLP architectures and hyperparameters to observe their effects on classification performance.

### Part 1: MLP on the Digits Dataset

1. Dataset: Use the `load\_digits` dataset from `sklearn.datasets`. This dataset contains images of handwritten digits (0–9), each represented as an 8x8 pixel grid (64 features).
2. Task: Build, train, and evaluate an MLP model to classify the digits.
3. Steps:

#### Data Preprocessing:

- Load the digits dataset using `sklearn.datasets.load\_digits()`.
- Scale the data using `StandardScaler` to normalize the feature values.

#### Model Architecture:

- Define an MLP model using Keras with the following structure:
- Input Layer: 64 features (input shape).
- Hidden Layers: At least two hidden layers, with 128 and 64 neurons, respectively, each with ReLU activation.
- Output Layer: 10 neurons with softmax activation (for the 10-digit classes).
- ```
mlp = MLPClassifier(hidden_layer_sizes=(128, 64), activation='relu', solver='adam',
max_iter=100)
mlp.fit(X_train, y_train)
```

#### Compile the Model:

- Use the Adam optimizer with a learning rate of 0.001.

- For the loss function, use categorical cross-entropy.

#### Training:

- Train the model for 30 epochs with a batch size of 16.

#### Evaluation:

- Evaluate the model on a test split (e.g., 20%) and report the test accuracy.

#### Visualization:

- Plot the training and validation accuracy over epochs to show model learning.
- Show the original images and model predictions on 5 samples from the test set.

#### 4. Hints:

- Use `train\_test\_split` from `sklearn.model\_selection` to split the data.
- Use `LabelBinarizer` to one-hot encode the target labels.

## Part 2: MLP on the Iris Dataset

1. Dataset: Use the `load\_iris` dataset from `sklearn.datasets`. This dataset consists of 150 samples of iris flowers, classified into three species based on four features (sepal and petal dimensions).

```
import numpy as np
import tensorflow as tf
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelBinarizer
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
import matplotlib.pyplot as plt
```

2. Task: Build, train, and evaluate a smaller MLP model to classify the iris species.

3. Steps:

#### **Data Preprocessing:**

- Load the Iris dataset using `sklearn.datasets.load\_iris()`.
- Standardize the data using `StandardScaler`.
- Binarize the data.

```
label_binarizer = LabelBinarizer()  
  
y_train = label_binarizer.fit_transform(y_train)  
  
y_test = label_binarizer.transform(y_test)
```

#### **Model Architecture:**

- Define an MLP model with the following structure:
- Input Layer: 4 features (input shape).
- Hidden Layers: Use one hidden layer with 64 neurons and ReLU activation.
- Output Layer: 3 neurons with softmax activation (for the three classes of iris).

```
# Define the MLP model  
  
model = Sequential([  
  
    Dense(64, input_shape=(X_train.shape[1],), activation='relu'), # First hidden  
    layer  
  
    Dense(32, activation='relu'), # Second hidden layer  
  
    Dense(3, activation='softmax') # Output layer with 3 classes (Iris Setosa,  
    Versicolor, Virginica)  
  
])
```

#### **Compile the Model:**

- Use the Adam optimizer with a learning rate of 0.01.
- Set the loss function to categorical cross-entropy.

```
# Compile the model  
  
model.compile(optimizer=Adam(learning_rate=0.01),  
              loss='categorical_crossentropy',
```

```
metrics=['accuracy'])
```

## Training:

- Train the model for 50 epochs with a batch size of 8.

```
# Train the model  
  
history = model.fit(X_train, y_train, epochs=50, batch_size=8,  
validation_split=0.1, verbose=1)
```

## Evaluation:

- Evaluate model accuracy on a test set (use a 20% test split) and report accuracy.

```
# Evaluate the model  
  
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=0)  
  
print(f"Test Accuracy: {test_accuracy:.2f}")
```

## Visualization:

- Plot the training and validation accuracy over epochs.
- Display a confusion matrix for the test predictions.

```
# Plot training and validation accuracy over epochs  
  
plt.plot(history.history['accuracy'], label='Training Accuracy')  
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.legend()  
plt.show()
```

## 4. Hints:

One-hot encode the labels with `LabelBinarizer`.

Use the `plot\_confusion\_matrix` function from `sklearn.metrics` to create the confusion matrix.

Experiment with different numbers of layers and neurons in and observe how model complexity impacts performance and training time.